

CircleUp Data Science Take-Home Assignment

March 31, 2018

1 Summary

We analysed the CircleUp dataset with python3 libraries *Pandas*, *Numpy*, *Scipy*, *Matplotlib*, *Seaborn* and *Scikit-learn*. Data have been locally downloaded in .csv files format and imported in Pandas data frames to be manipulated.

1.1 Question 1

To find the solutions we used the `user_message` dataset.

-

1.1.1 Point 1:

The list of users that added more than 500 pieces of content is the following:

User ID	Number of Content
9484	1163
9676	722
12116	688
3924	686
2052	640
10878	601
5999	566
8962	551
11578	544
2434	526
17616	526
3532	523
10530	521
11271	503

Details of the analysis are Section [1.3](#).

-

1.1.2 Point 2:

In order to find which users grew more rapidly in the last year, we need a measure of how fast they incremented the customer engagement in a given period. Such measurement can be obtained from a linear regression of the variable `total_engagement` versus the `content_created_date`. Then our metric (`growing_rate`), is represented by the slope of the fitted line. The bigger the slope, the faster the customer engagement growing. Finally, ranking the users based on this metric we found the top 10 fastest growing users:

user_id	growing_rate
3924	115.35
2052	91.11
4527	17.67
9484	9.31
12116	6.52
9770	4.33
3063	4.22
7215	3.61
5833	3.00
4711	2.73

Section 1.4 the details of the analysis, and more considerations about the manipulation of the variable `content_created_date`.

As last check, we added also some Section 1.5 of our `Growing_rate` metric.

1.2 Question 2

To answer the second question of the assignment, we build a supervised learning models using the `user`, `user_features` and `model_test_file` dataset. In order to predict the two classes of response outcome, we built and compared the performances of two different learning algorithms: **Support Vector Machine** and **Random Forest**.

As first, we performed an exploratory data analysis (Section 1.7) to better understand features properties and behaviour. We applied some features transformation as: eliminating isolated extreme values, taking logarithmic values and rescaling in order to understand possible correlations among features, and to put them in a better shape for feeding the model training.

We decided to train the classification models excluding few of the features inside the `user_features` dataset. We observed that the performance of the **SVM** learning algorithm is better than the **Random Forest**, especially in predicting the less represented *class* (response equal to 1). Details and results of the analysis are Section 1.6.

Looking for better results, we also used the complete list of features, conveniently transformed (see Section 1.12), to train the **SVM** model, but the Section 1.13 are completely comparable.]

Finally we predicted the response for the users in the `model_test_file`, using the **SVM** model trained with the subset of features. Section 1.14 the results.

1.3 Ranking of best users based on numbers of content created

```
In [3]: # Import libraries
```

```
import numpy as np
```

```
import pandas as pd
```

```
message_data = pd.read_csv('data/user_message.csv')
```

```
message_data.head()
```

```
In [4]: # Import dataset in Pandas data frame
```

```
message_data = pd.read_csv('data/user_message.csv')
```

```
message_data.head()
```

```
Out[4]:
```

	user_id	content_created_date	content_count	total_engagement
0	20	1/1/2015	1	52
1	20	1/2/2015	1	72
2	20	1/10/2015	1	83
3	20	1/12/2015	1	45
4	20	1/16/2015	1	102

```
In [5]: # Sum over content_count by user
```

```
df_by_user = message_data.groupby('user_id').sum()
```

```
In [6]: # Select and Sort
```

```
df = df_by_user[(df_by_user['content_count']>500)]
```

```
df.sort_values(by='content_count', ascending=False)
```

```
Out[6]:
```

	user_id	content_count	total_engagement
0	9484	1163	954382
1	9676	722	204249
2	12116	688	1243620
3	3924	686	16661950
4	2052	640	22826604
5	10878	601	194024
6	5999	566	373659
7	8962	551	53325
8	11578	544	90666
9	2434	526	30128
10	17616	526	180460
11	3532	523	91546
12	10530	521	16539
13	11271	503	71157

1.4 Fastest Growing users

In order to find the faster growing user, first of all, we need to reshape and transform our dataset. First step is to convert the variable `content_created_date` from *string* to *date*. Then, we transform each *date* in the corresponding *number of days since '2015-01-01'*.

```
In [7]: # Convert strings of date in a date object
message_data['content_created_date'] = pd.to_datetime(message_data['content_created_date'])
message_data.head()
```

```
Out[7]:
```

	user_id	content_created_date	content_count	total_engagement
0	20	2015-01-01	1	52
1	20	2015-01-02	1	72
2	20	2015-01-10	1	83
3	20	2015-01-12	1	45
4	20	2015-01-16	1	102

```
In [ ]: # Select only needed variables
df = message_data[['user_id', 'content_created_date', 'total_engagement']]
```

```
In [8]: # convert date in day_since to be able to apply regression
df['days_since'] = (df.content_created_date - pd.to_datetime('2015-01-01')).astype('timedelta64[D]')
df.head()
```

```
Out[8]:
```

	user_id	content_created_date	total_engagement	days_since
0	20	2015-01-01	52	0.0
1	20	2015-01-02	72	1.0
2	20	2015-01-10	83	9.0
3	20	2015-01-12	45	11.0
4	20	2015-01-16	102	15.0

To perform the linear regression with `stats.linregress`, we need to have for each user a list of `days_since`, and a list of `total_engagement` that correspond to our x and y variables.

```
In [11]: # reshape df to have a list of 'days_since' for each user
dfreshaped_x = df.groupby('user_id')['days_since'].apply(list)
dfreshaped_x.head()
```

```
Out[11]: user_id
20      [0.0, 1.0, 9.0, 11.0, 15.0, 21.0, 23.0, 25.0, ...
134     [28.0, 29.0, 32.0, 37.0, 62.0, 63.0, 64.0, 65...
635     [120.0, 121.0, 122.0, 123.0, 124.0, 125.0, 126...
950     [4.0, 49.0, 55.0, 56.0, 60.0, 61.0, 67.0, 75.0...
1034    [4.0, 11.0, 12.0, 13.0, 18.0, 20.0, 25.0, 41.0...
Name: days_since, dtype: object
```

```
In [12]: # reshape df to have a list of 'total_engagement' for each user
dfreshaped_y = df.groupby('user_id')['total_engagement'].apply(list)
dfreshaped_y.head()
```

```
Out[12]: user_id
20      [52, 72, 83, 45, 102, 35, 92, 33, 39, 109, 54, ...
134     [68, 44, 12, 13, 18, 69, 121, 23, 61, 34, 22, ...
635     [67, 37, 20, 36, 20, 7, 16, 25, 15, 4, 7, 14, ...
950     [18, 20, 8, 13, 6, 16, 8, 14, 12, 15, 15, 8, 1...
1034    [22, 28, 18, 12, 29, 26, 45, 50, 29, 29, 33, 2...
Name: total_engagement, dtype: object
```

```
In [16]: # Reshape data to have it in a useful format
data = {'days_since':dfreshaped_x, 'total_engagement':dfreshaped_y}
day_eng_list_df = pd.DataFrame(data)
day_eng_list_df.head()
```

```
Out[16]:
```

	days_since \
user_id	
20	[0.0, 1.0, 9.0, 11.0, 15.0, 21.0, 23.0, 25.0, ...
134	[28.0, 29.0, 32.0, 37.0, 62.0, 63.0, 64.0, 65...
635	[120.0, 121.0, 122.0, 123.0, 124.0, 125.0, 126...
950	[4.0, 49.0, 55.0, 56.0, 60.0, 61.0, 67.0, 75.0...
1034	[4.0, 11.0, 12.0, 13.0, 18.0, 20.0, 25.0, 41.0...

	total_engagement
user_id	
20	[52, 72, 83, 45, 102, 35, 92, 33, 39, 109, 54,...
134	[68, 44, 12, 13, 18, 69, 121, 23, 61, 34, 22, ...
635	[67, 37, 20, 36, 20, 7, 16, 25, 15, 4, 7, 14, ...
950	[18, 20, 8, 13, 6, 16, 8, 14, 12, 15, 15, 8, 1...
1034	[22, 28, 18, 12, 29, 26, 45, 50, 29, 29, 33, 2...

Another important analysis step is to check how many points (days_since, total_engagement) there are for each users. Because too few points could affect the quality of the linear regression fit resulting in an innacurate *slope* calculation. We choosed to exclude users with less than 10 entries in the data frame.

```
In [22]: # sanity check on the lenght of the lists in order to have a meaning regression points
day_eng_list_df['counts'] = day_eng_list_df['days_since'].apply(len)
day_eng_list_df.head()
```

```
Out[22]:
```

	days_since \	
user_id		
20	[0.0, 1.0, 9.0, 11.0, 15.0, 21.0, 23.0, 25.0, ...	
134	[28.0, 29.0, 32.0, 37.0, 62.0, 63.0, 64.0, 65...	
635	[120.0, 121.0, 122.0, 123.0, 124.0, 125.0, 126...	
950	[4.0, 49.0, 55.0, 56.0, 60.0, 61.0, 67.0, 75.0...	
1034	[4.0, 11.0, 12.0, 13.0, 18.0, 20.0, 25.0, 41.0...	

	total_engagement	counts
user_id		
20	[52, 72, 83, 45, 102, 35, 92, 33, 39, 109, 54,...	84
134	[68, 44, 12, 13, 18, 69, 121, 23, 61, 34, 22, ...	55
635	[67, 37, 20, 36, 20, 7, 16, 25, 15, 4, 7, 14, ...	126
950	[18, 20, 8, 13, 6, 16, 8, 14, 12, 15, 15, 8, 1...	113
1034	[22, 28, 18, 12, 29, 26, 45, 50, 29, 29, 33, 2...	73

```
In [23]: # Remove users with less than 10 entrie per list of days_since and total_engagement
cut_df = day_eng_list_df[day_eng_list_df['counts']>10]
```

Now, we define a function that performs the regression using `stats.linregress`, and returns the slope of the regression line.

```
In [25]: # Function definition
        from scipy import stats

        def linreg(row):
            slope, intercept, r_value, p_value, std_err = stats.linregress(row['days_since'], r
            return slope
```

Finally, we apply our `linreg` function to the reshaped data frame, and add a new column with the *slope* of the regression line. This new parameter is the `growing_rate` metric that we could use to find the faster growing users.

```
In [29]: # Apply linreg function to cut_df data frame
        cut_df['growing_rate'] = cut_df.apply(linreg, axis=1)
```

C:\Users\Michela\Anaconda3\lib\site-packages\ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>

The final answer to point 2 of Question 1 is the following table of the 10 users with the higher growing rate.

```
In [63]: # Sort and select first 10 users
        sorted_df = cut_df.sort_values('growing_rate', ascending=False)
        sorted_df[:10]
```

```
Out[63]:
```

	days_since \
user_id	
3924	[0.0, 1.0, 2.0, 4.0, 6.0, 7.0, 8.0, 10.0, 13.0...
2052	[0.0, 1.0, 2.0, 4.0, 5.0, 6.0, 7.0, 9.0, 11.0,...
4527	[0.0, 1.0, 6.0, 11.0, 16.0, 20.0, 23.0, 26.0, ...
9484	[109.0, 110.0, 111.0, 112.0, 113.0, 114.0, 115...
12116	[5.0, 6.0, 7.0, 8.0, 9.0, 10.0, 11.0, 12.0, 13...
9770	[17.0, 23.0, 53.0, 79.0, 80.0, 81.0, 82.0, 83...
3063	[4.0, 7.0, 11.0, 14.0, 18.0, 22.0, 23.0, 26.0,...
7215	[109.0, 110.0, 111.0, 112.0, 113.0, 114.0, 115...
5833	[4.0, 8.0, 13.0, 19.0, 22.0, 35.0, 36.0, 47.0,...
4711	[6.0, 9.0, 10.0, 12.0, 14.0, 16.0, 18.0, 20.0,...

	total_engagement	counts \
user_id		
3924	[21991, 56394, 26564, 17767, 36185, 21713, 236...	318
2052	[17273, 16608, 25792, 31585, 211897, 296452, 1...	298

4527	[274, 471, 308, 327, 362, 560, 374, 352, 295, ...	119
9484	[2039, 2994, 3229, 5582, 5287, 4302, 2554, 350...	198
12116	[1490, 4267, 5118, 907, 2507, 1038, 5417, 4488...	302
9770	[528, 10431, 4948, 4966, 4985, 529, 528, 488, ...	243
3063	[509, 370, 385, 574, 566, 538, 603, 422, 352, ...	81
7215	[136, 228, 164, 86, 340, 290, 180, 199, 196, 1...	219
5833	[2082, 1544, 1504, 688, 963, 1611, 928, 2973, ...	34
4711	[653, 143, 153, 124, 205, 163, 148, 229, 172, ...	185

	growing_rate
user_id	
3924	115.351638
2052	91.111386
4527	17.673102
9484	9.315306
12116	6.518727
9770	4.326168
3063	4.223012
7215	3.613537
5833	3.007240
4711	2.731612

1.5 Visualization of Growing rate

In order to check our results we visualized the regression line for the first and last ranked users.

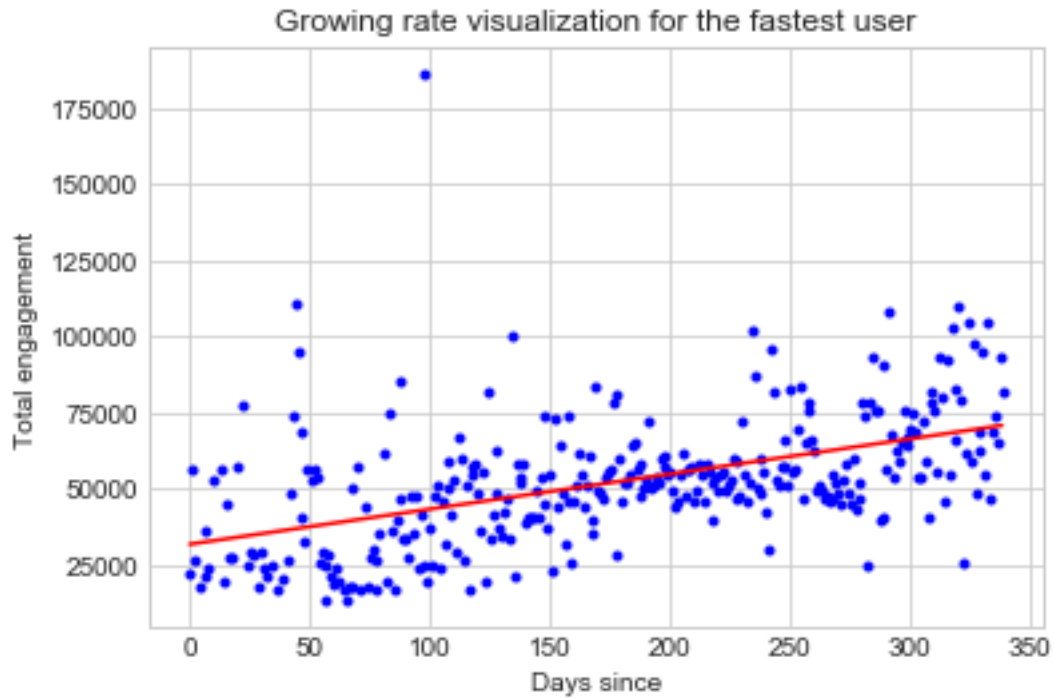
```
In [282]: # Select arrays of points for the first user and last users
x_coord=sorted_df.loc[:, 'days_since']
y_coord=sorted_df.loc[:, 'total_engagement']

x1 = x_coord.iloc[0]
y1 = y_coord.iloc[0]
x9 = x_coord.iloc[9]
y9 = y_coord.iloc[9]

In [288]: # Plot point and regression line
import matplotlib.pyplot as plt


plt.title('Growing rate visualization for the fastest user')
plt.xlabel('Days since')
plt.ylabel('Total engagement')
plt.plot(x1, y1, 'b.')
slope, intercept, r_value, p_value, std_err = stats.linregress(x1, y1)
plt.plot(x1, slope * np.array(x1) + intercept, 'r-')
print("slope =", slope)

slope = 115.3516377750555
```



```
In [287]: # Plot point and regression line
plt.title('Growing rate visualization for the last user')
plt.xlabel('Days since')
plt.ylabel('Total engagement')
plt.plot(x9, y9, 'b.')
slope, intercept, r_value, p_value, std_err = stats.linregress(x9, y9)
plt.plot(x9, slope * np.array(x9) + intercept, 'r-')
print("slope =", slope)
```

```
slope = 2.731612298914138
```

1.6 Learning Models

```
In [68]: # Import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [69]: # import data
features = pd.read_csv('data/user_features.csv')
users = pd.read_csv('data/user.csv')
```

```
In [70]: features.head()
```

```
Out[70]:
```

	var_1	var_2	var_3	var_4	var_5	var_6	var_7	var_8	var_9	\
0	1	1	0	1740115	38.637588	26	0	0.0000	0	
1	1	1	362	9803192	30.901522	11	237	1.7553	0	
2	1	1	3279	3694516	45.933998	329	638	13.1270	0	
3	1	1	213	18185084	34.120554	162	65	2.4769	0	
4	0	0	226	0	23.700552	3	99	5.0404	0	

	var_10	var_11	var_12	user_id
0	2542	0	0.87499	154531
1	117	17	0.79306	156315

2	1313	380	0.94425	149607
3	168	12	0.57251	26755
4	119	156	0.45718	149734

```
In [71]: users.head()
```

```
Out[71]:   response  user_id
0         0   154531
1         0   156315
2         0   149607
3         0    26755
4         0   149734
```

```
In [72]: # put togheter to have the final dataset
train = users.set_index('user_id').join(features.set_index('user_id'))
```

```
In [73]: train.head()
```

```
Out[73]:   response  var_1  var_2  var_3  var_4  var_5  var_6  var_7  \
user_id
154531         0      1      1      0  1740115  38.637588      26      0
156315         0      1      1     362  9803192  30.901522      11     237
149607         0      1      1    3279  3694516  45.933998     329     638
26755         0      1      1     213  18185084  34.120554     162     65
149734         0      0      0     226         0  23.700552       3     99


      var_8  var_9  var_10  var_11  var_12
user_id
154531  0.0000      0    2542       0  0.87499
156315  1.7553      0     117      17  0.79306
149607 13.1270      0    1313     380  0.94425
26755   2.4769      0     168      12  0.57251
149734  5.0404      0     119     156  0.45718
```

1.7 Exploratory Data Analysis

The first and most important step for the analysis is to try to have a clear understanding of our data. In order to do this, we checked for missing values, controlled the *balance* of the two sample of response, and took a closer look at the feautres.

```
In [22]: # Check for Missing Data
sns.heatmap(train.isnull(),yticklabels=False,cbar=False,cmap='viridis')
```

```
Out[22]: <matplotlib.axes._subplots.AxesSubplot at 0x21bf44bde80>
```




From the above heatmap is clear that the dataset does not contain any null value. All the values in the dataset are numeric, so no operation on the features are needed up to this point.

It is important to check which is the ratio between the two classes of response in the dataset, to control if it is balanced or not.

```
In [75]: # Chek the percentage of outcome
sns.set_style('whitegrid')
sns.countplot(x='response', data=train)
```

```
Out[75]: <matplotlib.axes._subplots.AxesSubplot at 0x23719355710>
```



```
In [76]: # Calculate the % of sample
         train['response'].value_counts()
```

```
Out[76]: 0    4516
         1     801
         Name: response, dtype: int64
```


The two classes of response are not equally represented. 82.3% of sample is 0 and only 17.7% is 1. This could affect the learning model. The predictive model could be biased and inaccurate because the second class is poor represented.

1.8 A closer look at the features

Vars_1 and var_2 are categorical variables (0,1), no continuous values. So we could take a look at their relation with the outcome response in order to check if they are strong predictor or not. In the *histograms* below we plotted the response variable using color to distinguish the values 0 or 1 in var_1 and var_2 variables respectively.

```
In [289]: # Check if var_1 is a strong predictor or not
          sns.set_style('whitegrid')
          sns.countplot(x='var_1', hue='response', data=train)
```

```
Out[289]: <matplotlib.axes._subplots.AxesSubplot at 0x2371ad653c8>
```




As you can see, almost all the values of `response = 1` are contained in `var_1 = 1`. In other words, we can be pretty confident that when `var_1` is 0, the probability to get a `response = 1` is very poor. This is certainly a property that our classifier will take advantage of.

A similar behaviour is observed in `var_2` (see plot below).

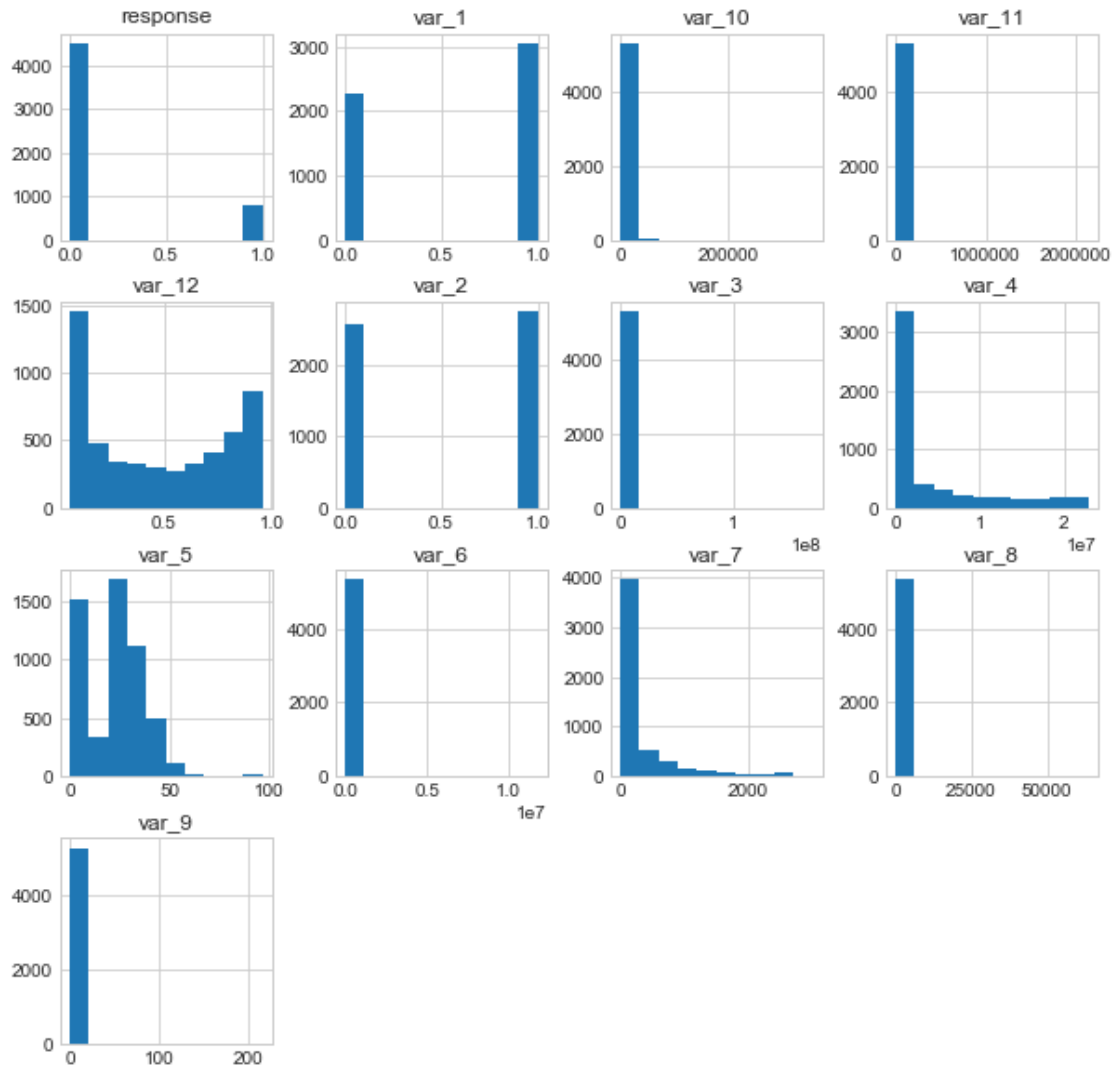
```
In [290]: # Check if var_2 is a strong predictor or not
sns.set_style('whitegrid')
sns.countplot(x='var_2', hue='response', data=train)
```

```
Out[290]: <matplotlib.axes._subplots.AxesSubplot at 0x23720517898>
```



Let's take a look to the other variables distribution and their possible correlations. In fact, not having self-explaining names, we do not know what they are and how they could be related to the outcome and between them.

```
In [80]: # Features plots  
train.hist(figsize=(10,10));
```



From the above distributions it is clear that variables var_3, var_4, var_6, var_8, var_9, var_10 and var_11 present lot of outliers and span multiple order of magnitude. So we could take the logarithmic value and then check for possible correlations.

1.9 Taking logarithmic value of some features


In order to apply the logarithmic transformation to misbehaving features we created a proper function: `log_transform`, that takes as input the dataset, and the list of variable to transform, and returns the modified data frame.

```
In [82]: # Define a function to apply log to a list of variables
def log_transform(df, var_list):
    for var in var_list:
        df[var] = df[var].apply(lambda x: np.log10(x+1))
```

```
In [83]: # apply the func
         log_transform(train, ['var_3','var_4','var_6','var_8','var_9','var_10','var_11'])
```

Now we could chek for possible correlations visually and than, after that, take a look at the Person coefficient to have a numerical estimation of the correlation.


```
In [84]: # check the correlation
         pd.plotting.scatter_matrix(train[['var_3','var_4','var_6','var_8','var_9','var_10','var_11']],
         plt.show())
```



From the above scatter plots it appear that with the exception of var_9, the other features seems to be all correlated with different intensity. For var_8 and var_11 the correlation is really strong. They are also correlated with var_3. We have checked also the Person coefficient to have a numeric estimation of the correlation.


```
In [85]: # check the pearsons coefficients
corr_matrix = train.corr()
#corr_matrix
# plot the heatmap
plt.subplots(figsize=(9,9))
sns.heatmap(corr_matrix, annot=True, xticklabels=corr_matrix.columns, yticklabels=corr_
```

Out[85]: <matplotlib.axes._subplots.AxesSubplot at 0x2371b9e95c0>



So, from the above heatmap appears also some correlation between *normal* features and *mis-behaving*. For example, var_5 and var_6, var_10 and var_2. Furthermore, var_4 seems to be an-

ticorrelated with `var_3`, and slightly also with `var_6` and `var_8`. So we decide to try to use only a subsample of features to train the learnig model.

1.10 Support Vector Machine Classifier (SVM)

As first attempt to model our dataset, we build a SVM classifier using only a subset of features. So first of all, we selected the features.

```
In [176]: # Subsetting the train dataset
trainmod = train[['response', 'var_1', 'var_2', 'var_5', 'var_7', 'var_12']]
trainmod.head()
```

```
Out[176]:
```

	response	var_1	var_2	var_5	var_7	var_12
user_id						
154531	0	1	1	38.637588	0	0.87499
156315	0	1	1	30.901522	237	0.79306
149607	0	1	1	45.933998	638	0.94425
26755	0	1	1	34.120554	65	0.57251
149734	0	0	0	23.700552	99	0.45718

Un important step in the building and training our model is to divide our dataset in a train and a test sample. So we could use the test sample to evaluate the model performances.

```
In [190]: # Split the dataset in train and test
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(trainmod.drop('response',axis=1),
                                                    trainmod['response'], test_size=0.
                                                    random_state=101)
```

1.10.1 Features Scaling

In order to have better results from the **SVM** classifier we need to rescale our features.

```
In [191]: # Features rescaling
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()

scaler.fit(X_train)
X_train=scaler.transform(X_train)
X_test=scaler.transform(X_test)
```

1.10.2 Build the classifier

Now we are ready to build the classifier. We used the `class_weight='balanced'` parameter to try to taking into account the bad balance of our dataset, where the *class* 0 is the most represented.

```
In [179]: # Buiding the classifier
from sklearn.svm import SVC
svc_model = SVC(class_weight='balanced')
```

```
In [180]: # Model fit
          svc_model.fit(X_train,y_train)

Out[180]: SVC(C=1.0, cache_size=200, class_weight='balanced', coef0=0.0,
            decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
            max_iter=-1, probability=False, random_state=None, shrinking=True,
            tol=0.001, verbose=False)

In [181]: # Model predictions
          svc_pred = svc_model.predict(X_test)
```

1.10.3 Model evaluation

Finally we need to evaluate our model in order to estimate his performances. We could take a look at the *Confusion Matrix* and at some estimators like *accuracy*, *precision*, *recall* and *F1-score*.

```
In [182]: # Model eveluation
          import itertools
          from sklearn.metrics import classification_report, confusion_matrix
          from sklearn import metrics

          def plot_confusion_matrix(cm, classes,
                                    normalize=False,
                                    title='Confusion matrix',
                                    cmap=plt.cm.Blues):
              """
              This function prints and plots the confusion matrix.
              Normalization can be applied by setting `normalize=True`.
              """
              if normalize:
                  cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
                  print("Normalized confusion matrix")
              else:
                  print('Confusion matrix, without normalization')

              print(cm)

              plt.imshow(cm, interpolation='nearest', cmap=cmap)
              plt.title(title)
              #plt.colorbar()
              tick_marks = np.arange(len(classes))
              #plt.xticks(tick_marks, classes, rotation=45)
              plt.xticks(tick_marks, classes)
              plt.yticks(tick_marks, classes)

              fmt = '.2f' if normalize else 'd'
              thresh = cm.max() / 2.
              for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
                  plt.text(j, i, format(cm[i, j], fmt),
```

```

        horizontalalignment="center",
        color="white" if cm[i, j] > thresh else "black")

plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')

# Compute confusion matrix
cnf_matrix = confusion_matrix(y_test,svc_pred)
np.set_printoptions(precision=2)

class_names = ['0', '1']

# Plot non-normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=class_names,
                      title='Confusion matrix, without normalization')

# Plot normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=class_names, normalize=True,
                      title='Normalized confusion matrix')

plt.show()
#accuracy = float(cnf_matrix[0][0] + cnf_matrix[1][1])/cnf_matrix.sum()
#print(accuracy)
print(classification_report(y_test,svc_pred))



```

Confusion matrix, without normalization

```
[[1043  281]
 [  43  229]]
```

Normalized confusion matrix

```
[[0.79 0.21]
 [0.16 0.84]]
```



	precision	recall	f1-score	support
0	0.96	0.79	0.87	1324
1	0.45	0.84	0.59	272
avg / total	0.87	0.80	0.82	1596

With the **SVM** classifier we obtained a **True Positive Rate** (*Recall* or *Sensitivity*) of 79% for response class 0 and 84% for response class 1. Those values are not so high, but correctly balanced, taking into account the different weight of the two samples.

1.11 Random Forest Classifier

In order to try to obtain a better classification for our dataset, we build also a **Random Forest classifier**, with the same subset of features.

```
In [95]: from sklearn.ensemble import RandomForestClassifier
         rfc = RandomForestClassifier(class_weight='balanced', n_estimators=101)

In [96]: rfc.fit(X_train, y_train)

Out[96]: RandomForestClassifier(bootstrap=True, class_weight='balanced',
                                criterion='gini', max_depth=None, max_features='auto',
                                max_leaf_nodes=None, min_impurity_decrease=0.0,
                                min_impurity_split=None, min_samples_leaf=1,
                                min_samples_split=2, min_weight_fraction_leaf=0.0,
                                n_estimators=101, n_jobs=1, oob_score=False, random_state=None,
                                verbose=0, warm_start=False)

In [97]: rfc_pred = rfc.predict(X_test)

In [117]: # Compute confusion matrix
          cnf_matrix = confusion_matrix(y_test, rfc_pred)
          np.set_printoptions(precision=2)

          class_names = ['0', '1']

          # Plot non-normalized confusion matrix
          plt.figure()
          plot_confusion_matrix(cnf_matrix, classes=class_names,
                               title='Confusion matrix, without normalization')

          # Plot normalized confusion matrix
          plt.figure()
```

```

plot_confusion_matrix(cnf_matrix, classes=class_names, normalize=True,
                      title='Normalized confusion matrix')

plt.show()
#accuracy = float(cnf_matrix[0][0] + cnf_matrix[1][1])/cnf_matrix.sum()
#print(accuracy)
print(classification_report(y_test,rfc_pred))

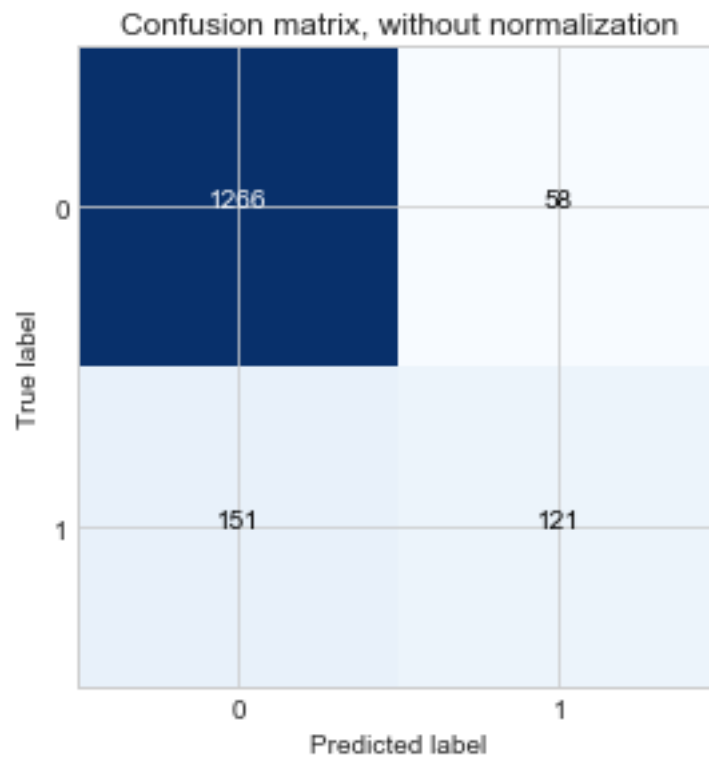
```

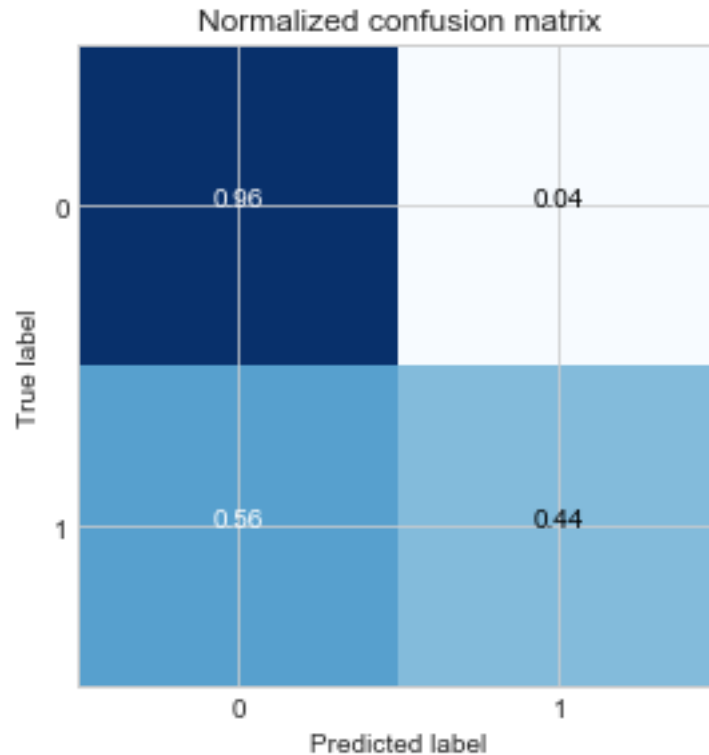
Confusion matrix, without normalization

```
[[1266  58]
 [ 151 121]]
```

Normalized confusion matrix

```
[[0.96 0.04]
 [0.56 0.44]]
```





	precision	recall	f1-score	support
0	0.89	0.96	0.92	1324
1	0.68	0.44	0.54	272
avg / total	0.86	0.87	0.86	1596

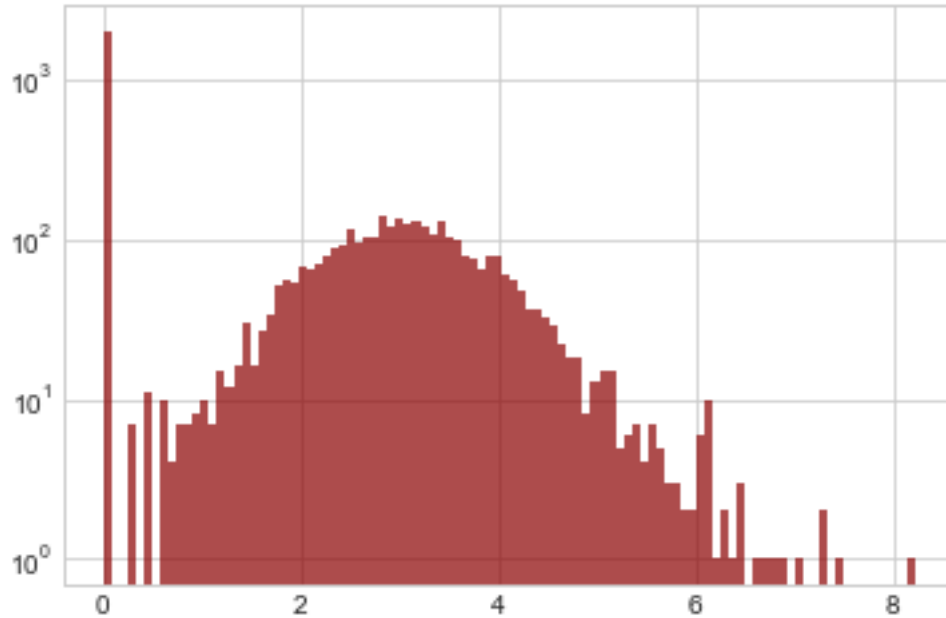
Using the **Random Forest** classifier the *sensitivity* for response 0 is quite better than the results of the **SVM**, respectively 96% and 79%. But for response 1 the *True Positive Rate* is very bad, 44% versus 84% with **SVM** classifier. So, we could say that the **Support Vector Machine** represents a good model to fit our dataset, and we will use it for *predictions* of new data.

1.12 Data Cleaning and Features Transformation

As last step of our analysis we tried to re-apply the learning model using all the features in order to check if the algorithm performances became better. But, to be useful, var_3, var_4, var_6, var_9 and var_10 need to be cleaned. Following plots show a closer look at those *critical* variables. We cutted the outliers values and took the logarithmic values in order to reduce the number of order of magnitude spanned from the variables.


```
In [118]: # Plot var_3
          train['var_3'].hist(bins=100,color='darkred',alpha=0.7, log=True)
```


Out[118]: <matplotlib.axes._subplots.AxesSubplot at 0x2371e3b8400>



In [119]: `train['var_4'].hist(bins=50,alpha=0.7,color='darkred',log=True)`


Out[119]: <matplotlib.axes._subplots.AxesSubplot at 0x23719a1c8d0>



It seems all the values of var_4 are concentrated at 0 and then there is a very long tail.


```
In [120]: train['var_6'].hist(bins=50,color='darkred',alpha=0.7, log=True)
```

```
Out[120]: <matplotlib.axes._subplots.AxesSubplot at 0x2371b22eeb8>
```




```
In [121]: train['var_9'].hist(bins=50,color='darkred',alpha=0.7, log=True)
```

```
Out[121]: <matplotlib.axes._subplots.AxesSubplot at 0x2371e1376d8>
```



```
In [122]: train['var_10'].hist(bins=50,color='darkred',alpha=0.7, log=True)
```

```
Out[122]: <matplotlib.axes._subplots.AxesSubplot at 0x2371b3ebba8>
```



1.13 Training SVM model after data cleaning

After features transformation, we applied the **SVM** classifier with the complete list of features, removing only features highly correlated: var_4, var_8 and var_11.

```
In [233]: # Subsetting the train dataset
          new_trainmod = train[['response', 'var_1', 'var_2', 'var_3', 'var_5', 'var_6', 'var_7', 'v
```

```
In [234]: # Split samples
          X_train, X_test, y_train, y_test = train_test_split(new_trainmod.drop('response', axis=
                                                                new_trainmod['response'], test_size=
                                                                random_state=101)
```

```
In [235]: # Rescaling features
          newscaler = StandardScaler()
          newscaler.fit(X_train)
          X_train = newscaler.transform(X_train)
          X_test = newscaler.transform(X_test)
```

```
In [236]: # Model fit
          svc_new_model = SVC(class_weight='balanced')
          svc_new_model.fit(X_train, y_train)
```

```

Out[236]: SVC(C=1.0, cache_size=200, class_weight='balanced', coef0=0.0,
            decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
            max_iter=-1, probability=False, random_state=None, shrinking=True,
            tol=0.001, verbose=False)

In [237]: # Model predictions
          svc_new_pred = svc_new_model.predict(X_test)

In [238]: # Compute confusion matrix
          cnf_matrix = confusion_matrix(y_test, svc_new_pred)
          np.set_printoptions(precision=2)

          class_names = ['0', '1']

          # Plot non-normalized confusion matrix
          plt.figure()
          plot_confusion_matrix(cnf_matrix, classes=class_names,
                                title='Confusion matrix, without normalization')

          # Plot normalized confusion matrix
          plt.figure()
          plot_confusion_matrix(cnf_matrix, classes=class_names, normalize=True,
                                title='Normalized confusion matrix')

          plt.show()
          #accuracy = float(cnf_matrix[0][0] + cnf_matrix[1][1])/cnf_matrix.sum()
          #print(accuracy)
          print(classification_report(y_test,svc_new_pred))



```

Confusion matrix, without normalization

```
[[1046  278]
 [  44 228]]
```

Normalized confusion matrix

```
[[0.79 0.21]
 [0.16 0.84]]
```



	precision	recall	f1-score	support
0	0.96	0.79	0.87	1324
1	0.45	0.84	0.59	272
avg / total	0.87	0.80	0.82	1596

New results are completely comparable with the model builded and trained with only the subset of the features. The inclusion of the other features did not improved the performances of the learning algorithm, as expected. Indeed, even after the *logarithmic* transformation most part of the features values are concentrated at zero, suppressing the weight of the other values. Unfortunately we cannot remove those *sparks* at zero from the dataset without loosing the most part of it. We tried and we obtained only few hundreds of data, quite unuseful to build and train whatever learning model.

In search of better performances, we also used **GridSearch** to try to optimize the paramenters *C* and *gamma* of the **SVM** models, but with the option *class_weight='balanced'* GridSearchCV diverged resulting in a totally wrong classification.

1.14 New data classification with SVM algorithm

At this point we will use the trained **SVM** model to predict the response on a new dataset. First of all, we need to import the new data in a Pandas dataframe and apply all the transformation we did on training sample.

```
In [199]: # Import data
test = pd.read_csv('data/model_test_file.csv')
test.head()
```

```
Out[199]:
```

	var_1	var_2	var_3	var_4	var_5	var_6	var_7	var_8	var_9	\
0	1	1	193	11248797	33.254023	17	40	77.7500	0	
1	1	1	14631	2396429	42.681445	1005	1582	56.3211	0	
2	1	1	385	3521953	33.501102	294	105	8.4952	26	
3	0	0	0	0	0.000000	0	0	0.0000	0	
4	1	1	6672	2913930	34.344784	10582	30	0.7667	0	

	var_10	var_11	var_12	user_id
0	583	3055	0.60563	154472
1	4082	53187	0.88451	151147
2	162	37	0.86154	10543
3	0	0	0.10018	136986
4	133	7	0.68598	137008

We need to remove same features as in the training analysis that results in a better performances. So we exclude *var_3*, *var_4*, *var_6*, *var_8*, *var_9*, *var_10* and *var_11*.

```
In [200]: # Drop misbehaving features
#test.drop(['var_3', 'var_4', 'var_6', 'var_8', 'var_9', 'var_10', 'var_11', 'user_id'], axis=
testmod = test.drop(['var_3', 'var_4', 'var_6', 'var_8', 'var_9', 'var_10', 'var_11'], axis=
# Set user_id as index
testmod = testmod.set_index('user_id')
testmod.head()
```

```
Out[200]:
```

	var_1	var_2	var_5	var_7	var_12
user_id					
154472	1	1	33.254023	40	0.60563
151147	1	1	42.681445	1582	0.88451
10543	1	1	33.501102	105	0.86154
136986	0	0	0.000000	0	0.10018
137008	1	1	34.344784	30	0.68598

Now, we need to *reshape* a bit the data frame to have the right format for the learning model.

```
In [202]: # Reshaping
X_test_newsample = testmod.as_matrix()
```

Now we are ready to apply the **SVM** model we builded and trained with only the *good* features. But first we have to *rescale* also the new test data.

```
In [203]: # rescaling
X_test_newsample = scaler.transform(X_test_newsample)
```

The predictions of our **SVM** learning model are:

```
In [204]: # apply SVM
svc_new_pred = svc_model.predict(X_test_newsample)
```

```
In [208]: # see the results taking a look at the predictions created
testmod['new_pred'] = svc_new_pred
testmod
```

```
Out[208]:
```

	var_1	var_2	var_5	var_7	var_12	new_pred
user_id						
154472	1	1	33.254023	40	0.60563	0
151147	1	1	42.681445	1582	0.88451	1
10543	1	1	33.501102	105	0.86154	1
136986	0	0	0.000000	0	0.10018	0
137008	1	1	34.344784	30	0.68598	0
137059	0	0	0.000000	0	0.08061	0
137253	0	0	0.000000	0	0.08061	0
137323	1	1	29.925298	20	0.77864	1
171652	1	1	0.000000	1810	0.67349	0
137341	1	1	34.572958	399	0.83240	1
137768	1	1	22.663224	236	0.25554	0
137822	1	1	32.954738	450	0.59736	0

137879	0	0	26.804861	217	0.71104	0
137984	1	1	35.802761	190	0.92892	1
138103	1	1	0.000000	1	0.17138	0
138200	1	1	26.361832	380	0.55178	0
138405	1	1	38.972069	563	0.92875	1
138439	0	0	15.813438	0	0.38118	0
138473	0	0	22.014121	0	0.07780	0
138483	0	0	26.903693	0	0.65279	0
138582	0	0	0.000000	0	0.08061	0
138606	1	0	0.000000	85	0.15566	0
138738	0	0	21.820799	240	0.10255	0
138948	0	0	0.000000	0	0.08061	0
139606	0	0	0.000000	0	0.14191	0
149425	1	1	43.409466	1563	0.93951	1
149419	1	0	27.236982	36	0.81458	1
150894	1	1	51.120647	1342	0.89298	1
141733	0	0	18.876661	0	0.20673	0
144653	1	1	24.671342	222	0.90147	1
147263	1	1	39.009098	27	0.81860	1
147793	0	0	0.000000	0	0.10714	0
148663	1	1	43.537221	957	0.93709	1
149065	1	0	29.102700	3	0.33104	0
149643	0	0	0.000000	0	0.41157	0
149792	1	1	0.000000	835	0.16774	0
10420	1	0	48.228613	0	0.35088	0
152535	1	1	22.981221	69	0.38972	0
153169	1	1	0.000000	457	0.73181	1
153389	0	1	0.000000	0	0.11321	0
153395	1	0	0.000000	55	0.35629	0
154070	0	0	0.000000	0	0.08061	0
154777	0	0	0.000000	0	0.08061	0
168361	0	0	0.000000	1123	0.10630	0
168689	1	1	26.400655	695	0.31953	0
168840	1	0	26.877454	5	0.50921	0
168845	1	1	0.000000	17	0.28022	0
168862	1	1	28.673564	1403	0.65822	0
168864	1	1	21.923702	13	0.69437	0
168989	1	1	0.000000	183	0.29884	0
168991	1	1	30.403461	702	0.75216	1
169591	0	0	0.000000	0	0.06707	0
169636	1	0	0.000000	4	0.10986	0
169853	0	0	23.274683	0	0.30361	0
259509	1	1	26.417065	787	0.80725	1
260608	1	1	40.801757	885	0.94547	1
602671	1	1	24.499112	101	0.52842	0
602672	0	0	15.846660	0	0.09382	0
8107	1	1	45.499149	89	0.77027	1